

# Flow GLM

XIN Baiying

September 25, 2025

## Notations

---

- $Y \in \mathbb{N}_0^{N \times G}$ : observed counts, N samples, G genes. For each cell  $i$ :
- Perturb label  $t_i \in \{\text{non-targeting}\} \cup \mathcal{G}_{\text{pert}}$ , where  $\mathcal{G}_{\text{pert}}$  is the set of targeting perturbations.
- Batch label  $b_i \in \{1, \dots, B\}$ .

# ZINB/Hurdle Modelling I

---

We still try to formulate the observation  $Y_{ig}$  as some count distribution. Here, we consider Zero-Inflated Negative Binomial (ZINB) distribution and Hurdle distribution.

**For ZINB:** For each gene  $g$ , assume  $Y_{ig}$  follows a ZINB distribution with mean  $\mu_{ig} > 0$ , dispersion  $\theta_g > 0$  and zero-inflation probability  $\pi_g \in [0, 1]$ :

$$\Pr(Y_{ig} = 0) = \pi_g + (1 - \pi_g) \cdot \Pr_{\text{NB}}(0; \mu_{ig}, \theta_g)$$

$$\Pr(Y_{ig} = y > 0) = (1 - \pi_g) \cdot \Pr_{\text{NB}}(y; \mu_{ig}, \theta_g).$$

- Negative Binomial (NB) can be constructed as a Gamma-Poisson mixture:  
 $\lambda_{ig} \sim \text{Gamma}(\theta_g, \frac{\theta_g}{\mu_{ig}})$ ,  $Y_{ig} | \lambda_{ig} \sim \text{Poisson}(\lambda_{ig})$ . Thus, for inference, we first let  $Y_{ig} = 0$  with probability  $\pi_g$ , otherwise first sample  $\lambda_{ig} \sim \text{Gamma}(\theta_g, \frac{\theta_g}{\mu_{ig}})$  then  $Y_{ig} | \lambda_{ig} \sim \text{Poisson}(\lambda_{ig})$ .

## ZINB/Hurdle Modelling II

---

**For Hurdle:** Quite similar to ZINB, except that we separately model the zero and non-zero counts:

$$\Pr(Y_{ig} = 0) = 1 - \rho_{ig},$$

$$\Pr(Y_{ig} = y > 0) = \rho_{ig} \cdot \frac{\Pr_{\text{NB}}(y; \mu_{ig}, \theta_g)}{1 - \Pr_{\text{NB}}(0; \mu_{ig}, \theta_g)}.$$

- Here  $\rho_{ig} = \text{sigmoid}(\rho_{0g} + \alpha u_{ig})$ , where  $\rho_{0g} \in \mathbb{R}$  is a learnable intercept,  $u_{ig}$  is the linear predictor (to be defined later), and  $\alpha > 0$  is also a learnable parameter to control the effect of  $u_{ig}$  on  $\rho_{ig}$ .

## Flow-GLM: Architecture I

---

Something different is that here we let  $\mu_{ig}$  to be obtained from a **Flow-GLM** model:

$$\mu_{ig} = s_g f(u_{ig}) s_i,$$

- $f : \mathbb{R} \rightarrow (0, \infty)$  is a strictly monotone (thus invertible) link function, which is called **flow function**. The details will be discussed later.
- $u_{ig}$  is the linear predictor for cell  $i$  and gene  $g$ :

$$u_{ig} = \log \mu_g^{\text{ctrl}} + \beta_g + (GKP(t_i))_g + h_{b_i, g}, \quad \forall i, g$$

- $\mu_g^{\text{ctrl}}$  is the baseline mean expression for gene  $g$  in control cells.
- $\beta_g \in \mathbb{R}$  is a learnable intercept shift for gene  $g$ .
- $(GKP(t_i))_g$  is basically the same as our previous NB model.
  - $G \in \mathbb{R}^{G \times d}$ , gene embedding, is obtained by: first construct the pseudo-bulk PB  $\in \mathbb{R}^{|\mathcal{G}_{\text{pert}}| \times G}$  by aggregating the counts of all cells with the same perturbation; then do PCA on  $\text{PB}^\top$  and take the first  $d$  principal components, finally normalize.

## Flow-GLM: Architecture II

---

- $P(t_i) \in \mathbb{R}^d$ , perturbation embedding of  $t_i$ , is obtained by  $P(t_i) = W_{\text{pert}} G_{t_i}$ , where  $W_{\text{pert}} \in \mathbb{R}^{d \times d}$  is a learnable weight matrix and  $G_{t_i} \in \mathbb{R}^d$  is the gene embedding of the perturbed gene  $t_i$  from the matrix  $G$ .
- $K \in \mathbb{R}^{d \times d}$  is a learnable linear transformation matrix.
- $h_{b_i, g}$  is the learnable batch effect for batch  $b_i$  and gene  $g$ . Basically, we assign each batch  $b_i$  an embedding vector  $e_{b_i} \in \mathbb{R}^r$  and then use a learnable weight matrix  $W_{\text{batch}} \in \mathbb{R}^{r \times G}$  to get  $h_{b_i, g} = (e_{b_i}^\top W_{\text{batch}})_g$ .
- $s_i > 0$  is the size factor for cell  $i$ , calculated as  $s_i = \frac{\sum_g Y_{ig}}{\text{medianUMI}_{\text{ctrl}}}$
- $s_g > 0$  is a learnable gene-specific scaling factor. In practice, we set  $s_g = \text{softplus}(\gamma_g) + \epsilon$ , where  $\gamma_g \in \mathbb{R}$  is a learnable parameter and  $0 < \epsilon \ll 1$  is a small constant to avoid numerical issues. Moreover, we may set a  $\ell_2$  regularization on  $\log s_g$  (or equivalently on  $\gamma_g$ ) to restrict the scale of  $s_g$ .

Recall  $\boxed{\mu_{ig} = s_g f(u_{ig}) s_i}$

## Flow Function

---

$f$  is required to be strictly monotone and map  $\mathbb{R}$  to  $(0, \infty)$ . Here, preliminarily we consider:

$$f(u) = \text{softplus} \left( a_0 + a_1 u + \sum_{k=1}^K w_k \text{softplus} (b_k(u - c_k)) \right), \quad a_1, w_k, b_k \geq 0$$

- $K$  is a hyperparameter to control the flexibility of  $f$ .
- $a_0, a_1, w_k, b_k, c_k$  are learnable parameters.

As a matter of fact,  $f$  can be further extended to more complex forms, e.g. RQ-Spline, or even neural networks with monotonicity constraints. We will explore these options in the future.

## Flow-GLM: Loss Function I

---

Generally, the loss function is composed of the following parts:

$$\mathcal{L} = \mathcal{L}_{\text{NLL}} + \lambda_{\text{bulk}} \mathcal{L}_{\text{bulk}} + \mathcal{L}_\theta + \mathcal{L}_\pi$$

**Negative log-likelihood loss:** For a batch of cells  $\mathcal{B}$ , recall that for ZINB,

$$\log p_{\text{ZINB}}(y; \mu, \theta, \pi) = \begin{cases} \log\left(\pi + (1 - \pi)p_{\text{NB}}(0; \mu, \theta)\right), & y = 0, \\ \log(1 - \pi) + \log p_{\text{NB}}(y; \mu, \theta), & y > 0, \end{cases}$$

where  $\log p_{\text{NB}}(y; \mu, \theta) = \log \Gamma(y + \theta) - \log \Gamma(\theta) - \log \Gamma(y + 1) + \theta \log \frac{\theta}{\theta + \mu} + y \log \frac{\mu}{\theta + \mu}$ . Collectively, the negative log-likelihood loss is

## Flow-GLM: Loss Function II

---

$$\mathcal{L}_{\text{NLL}} = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \sum_{g=1}^G \left( -\log p_{\text{ZINB}}(Y_{ig} \mid \mu_{ig}, \theta_g, \pi_g) \right)$$

- Here  $\mu_{ig} = s_g f(u_{ig}) s_i$  is obtained from the Flow-GLM model mentioned before.
- $\theta_g > 0$  and  $\pi_g \in [0, 1]$  can be directly estimated from the data using methods like method of moments, and can optionally choose to be further optimized during training by setting  $\ell_2$  regularization losses  $\mathcal{L}_\theta$  and  $\mathcal{L}_\pi$ .

The loss function for Hurdle model is quite similar, just replace  $p_{\text{ZINB}}$  with  $p_{\text{Hurdle}}$ .

## Flow-GLM: Loss Function III

---

**Pseudo-bulk loss:** This is a subsidiary loss to adjust the model to better fit the pseudo-bulk data in order to improve Perturbation Discrimination Score.

- In real training data, for each perturbation  $p \in \mathcal{G}_{\text{pert}}$ , calculate the pseudo-bulk average:  $\bar{\mathbf{y}}_p = \left[ \frac{1}{|\mathcal{C}_p|} \sum_{i \in \mathcal{C}_p} Y_{ig} \right]^\top \in \mathbb{R}^G$ , where  $\mathcal{C}_p = \{i : t_i = p\}$  is the set of cells with perturbation  $p$ .
- In ZINB model's prediction, the expected value of  $Y_{ig}$  is  $\mathbb{E}[Y_{ig}] = (1 - \pi_g)\mu_{ig}$ . Thus, we can also generate a pseudo-bulk average prediction for each perturbation  $p$  as  $\hat{\mu}_p = \left[ (1 - \pi_g) \frac{1}{|\mathcal{C}_p|} \sum_{i \in \mathcal{C}_p} \mu_{ig} \right]^\top \in \mathbb{R}^G$ .

## Flow-GLM: Loss Function IV

---

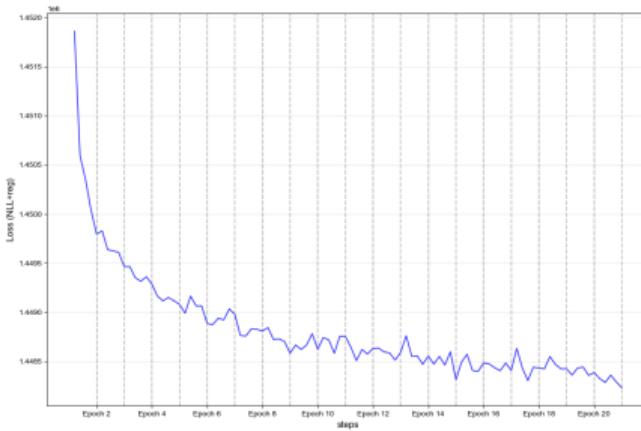
Then, we here use Huber loss to measure the difference between  $\bar{y}_p$  and  $\hat{\mu}_p$ , and try to minimize it:

$$\mathcal{L}_{\text{bulk}} = \frac{1}{|\tilde{\mathcal{G}}_{\text{pert}}|} \sum_{p \in \tilde{\mathcal{G}}_{\text{pert}}} \text{Huber}_{\delta}(\log(1 + \hat{\mu}_{pg}) - \log(1 + \bar{y}_{pg}))$$

- $\text{Huber}_{\delta}(r) = \begin{cases} \frac{1}{2}r^2, & |r| \leq \delta, \\ \delta(|r| - \frac{1}{2}\delta), & |r| > \delta. \end{cases}$  can be regarded as a combination of  $\ell_2$  loss and  $\ell_1$  loss, which is less sensitive to outliers than  $\ell_2$  loss.
- $\delta > 0$  is a hyperparameter to control the threshold between  $\ell_2$  and  $\ell_1$  loss, which can be tuned based on validation performance.
- Here  $\tilde{\mathcal{G}}_{\text{pert}} \subseteq \mathcal{G}_{\text{pert}}$  is a subset of perturbations ( $|\tilde{\mathcal{G}}_{\text{pert}}| \approx 2000$ ) that we randomly sampled as an estimation of the full set  $\mathcal{G}_{\text{pert}}$  to reduce computation cost.

# Experimental Results

---



- Differential Expression Score: 0.1746, Perturbation Discrimination Score: 0.5076, Mean Absolute Error: 0.2012, Overall Score: 2.6.
- Currently still unsatisfying, but it's still promising with a lot of rooms for improvement.

Aa	Name	Description	# Differen...	# Perturb...	# MAE	# Overall Score	Date
	📄 N 📁 OPEN 922	flow 5epoch	0.127	0.52	0.433	1.1	September 22, 2025
	📄 NB_flow_0 922	flow 10 epoch	0.123	0.509	0.336	0.7	September 22, 2025
	📄 NB_flow_0 922	flow, 10 epoch, reorganize	0.244	0.508	0.336	5.2	September 22, 2025
		test 50epoch	0.138	0.508	1.491	1.2	September 23, 2025
		test 50epoch_count_ robust	0.148	0.508	1.491	1.6	September 23, 2025
		test 10epoch_count_ fix	0.135	0.517	0.151	1.2	
		test 10epochnew	0.152	0.512	0.278	1.7	September 24, 2025
		test 20epoch plus new	0.168	0.51	0.39	2.3	September 24, 2025
		python flow_glm_old.py \ --train_h5ad /root/autodl- tmp/vcc_data/ad ata_Training.h5a d \ --epochs 10 \ /	0.249	0.508	0.339	5.3	September 24, 2025
		same but no count	0.124	0.502	0.339	0.7	
		gpu_ver	0.1746	0.5076	0.2012	2.6	