

# **TOSICA, GET, Baseline and Cell-eval**

XIN Baiying

August 26, 2025

# Overview

---

- 1. Transformer for one stop interpretable cell type annotation (TOSICA)**
- 2. A foundation model of transcription across human cell types (GET)**
- 3. VCC Baseline and Cell-eval**

# **Transformer for one stop interpretable cell type annotation (TOSICA)**

Chen J, Xu H, Tao W, Chen Z, Zhao Y, Han JJ

*Nature Communications*, 2023

DOI: 10.1038/s41467-023-35923-4

<https://www.nature.com/articles/s41467-023-35923-4.pdf>

# Introduction

---

The general goal of TOSICA is **cell type annotation**.

- Its input is a single-cell transcriptome expression profile  $\mathbf{e} \in \mathbb{R}^n$ , where  $n$  is the number of genes (10000+).
- The output is the predicted cell type  $\hat{y} \in \{1, \dots, C\}$ , where  $C$  is the number of cell types (10 ~ 50).

# TOSICA Architecture

---

TOSICA is based on the Transformer architecture and introduces a mask mechanism with biological prior knowledge. Its design consists of three layers:

- **Cell Embedding:** Maps each cell's expression to a low-dimensional representation space, where each dimension corresponds to a biological pathway/regulatory module.
- **Multi-Head Self-Attention:** Uses attention mechanism to compute attention between the cell embedding and classification CLS token to obtain contextual information for cell types.
- **Cell-Type Classifier:** Maps attention information to cell type space through fully connected layers to achieve final classification prediction.

## Cell Embedding I

---

For each cell  $\mathbf{e} \in \mathbb{R}^n$ , where  $n$  is the number of genes, we aim to map it to a low-dimensional representation through a fully connected network:

$$\mathbf{t} = \mathbf{W}'\mathbf{e} \in \mathbb{R}^k$$

where  $\mathbf{W}' \in \mathbb{R}^{k \times n}$  is a learnable weight matrix. In this  $k$ -dimensional space, each dimension serves as a token, and each token represents a biological pathway/regulatory module.

Since we want  $\mathbf{t}$  to have each dimension represent a specific biologically meaningful pathway in a structured manner, the weight matrix  $\mathbf{W}'$  needs to be specially designed with prior knowledge.

To integrate biological prior knowledge into the cell embedding process:

- First introduce a general weight matrix  $\mathbf{W} \in \mathbb{R}^{k \times n}$

## Cell Embedding II

---

- Based on prior knowledge, we construct a mask matrix  $\mathbf{M} \in \{0, 1\}^{k \times n}$ 
  - Where  $M_{ij} = 1$  if and only if the  $i$ -th pathway contains the  $j$ -th gene. This knowledge is obtained from external databases (specifically, gene set datasets from the Gene Set Enrichment Analysis database)
  - Therefore, only genes within the same pathway are mapped to the same token. This avoids the non-interpretability of mixed components, though it also limits the model's expressive power.
- Then perform Hadamard product (element-wise product) on these two matrices to obtain the final weight matrix:

$$\mathbf{W}' = \mathbf{M} \odot \mathbf{W}$$

$$\mathbf{t} = \mathbf{W}' \mathbf{e} \in \mathbb{R}^k$$

## Cell Embedding III

---

To enhance performance, the above operation is performed in parallel  $m$  times ( $m = 48$ ), concatenated to obtain:

$$\mathbf{T} := [\mathbf{t}_1 \quad \mathbf{t}_2 \quad \cdots \quad \mathbf{t}_m] \in \mathbb{R}^{k \times m}.$$

## Multi-Head Self-Attention I

---

First, we introduce a learnable dummy token  $\mathbf{cls} \in \mathbb{R}^m$  and concatenate it as the first row of  $\mathbf{T}$  to obtain the new input:

$$\mathbf{I} := \begin{bmatrix} \mathbf{cls}^\top \\ \mathbf{T} \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & \cdots & c_m \\ \mathbf{t}_1 & \mathbf{t}_2 & \cdots & \mathbf{t}_m \end{bmatrix} \in \mathbb{R}^{(1+k) \times m}.$$

- This is a common practice in Transformer and other NLP models. CLS is an additional, learnable virtual token. Since attention operations output a representation with the same shape as the input, we can use the output of this CLS as the global information extraction result for the current input, similar to a statistic of the current input, serving as the token for subsequent classification.

We feed this input  $\mathbf{I}$  into the attention mechanism. We first discuss the case of a single head. Overall, we have  $\mathbf{O} = \text{Attention}(\mathbf{I}) \in \mathbb{R}^{(1+k) \times m}$ . The specific computational details are as follows:

## Multi-Head Self-Attention II

---

- Through three linear transformations  $\mathbf{W}_q, \mathbf{W}_k, \mathbf{W}_v \in \mathbb{R}^{(1+k) \times (1+k)}$ , we map the input to Query, Key, and Value spaces:

$$\mathbf{Q} = \mathbf{W}_q \mathbf{I} \in \mathbb{R}^{(1+k) \times m},$$

$$\mathbf{K} = \mathbf{W}_k \mathbf{I} \in \mathbb{R}^{(1+k) \times m},$$

$$\mathbf{V} = \mathbf{W}_v \mathbf{I} \in \mathbb{R}^{(1+k) \times m}.$$

- We compute attention scores to obtain contextual information:

$$\mathbf{A} = \text{softmax} \left( \frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}} \right) \in \mathbb{R}^{(1+k) \times (1+k)}.$$

where  $d_k = m$  is the dimension of the Key vectors. This attention score matrix  $\mathbf{A}$  represents the similarity between different tokens.

## Multi-Head Self-Attention III

---

- We obtain contextual representations for each token through weighted summation:

$$\mathbf{O} = \mathbf{AV} \in \mathbb{R}^{(1+k) \times m}.$$

The multi-head case is similar to the above, equivalent to performing  $H$  computations in parallel.

- For the  $h$ -th head, we independently introduce learnable weight matrices  $\mathbf{W}_q^h, \mathbf{W}_k^h, \mathbf{W}_v^h \in \mathbb{R}^{(1+k) \times (1+k)}$ . We finally obtain the output  $\mathbf{O}^h = \text{Attention}^h(\mathbf{I}) \in \mathbb{R}^{(1+k) \times m}$ .
- We concatenate the  $H$  independent outputs and apply an additional linear mapping to reshape back to the original input shape:  
$$\tilde{\mathbf{O}} = \mathbf{W}_o [\mathbf{O}^1 \quad \mathbf{O}^2 \quad \dots \quad \mathbf{O}^H] \in \mathbb{R}^{(1+k) \times m}.$$

## Cell-Type Classifier

---

Finally, since the attention mechanism's input and output have the same shape, we extract the first row of  $\widetilde{\mathbf{O}}$  (denoted as  $\widetilde{\mathbf{cls}}$ ), which is the output of the input CLS token after the attention mechanism, as the global information extraction result for the current input.

We pass this result through a fully connected neural network to complete classification:

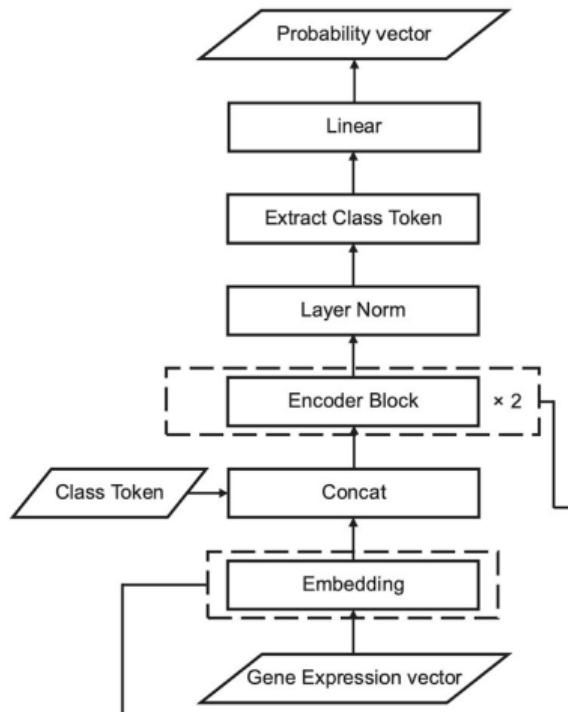
$$\mathbf{p} = \text{softmax}(\mathbf{W}_c \widetilde{\mathbf{cls}}) \in \mathbb{R}^C$$

where  $\mathbf{W}_c \in \mathbb{R}^{C \times m}$  is the learnable classification weight matrix, and  $C$  is the number of cell types.

# Model Architecture I

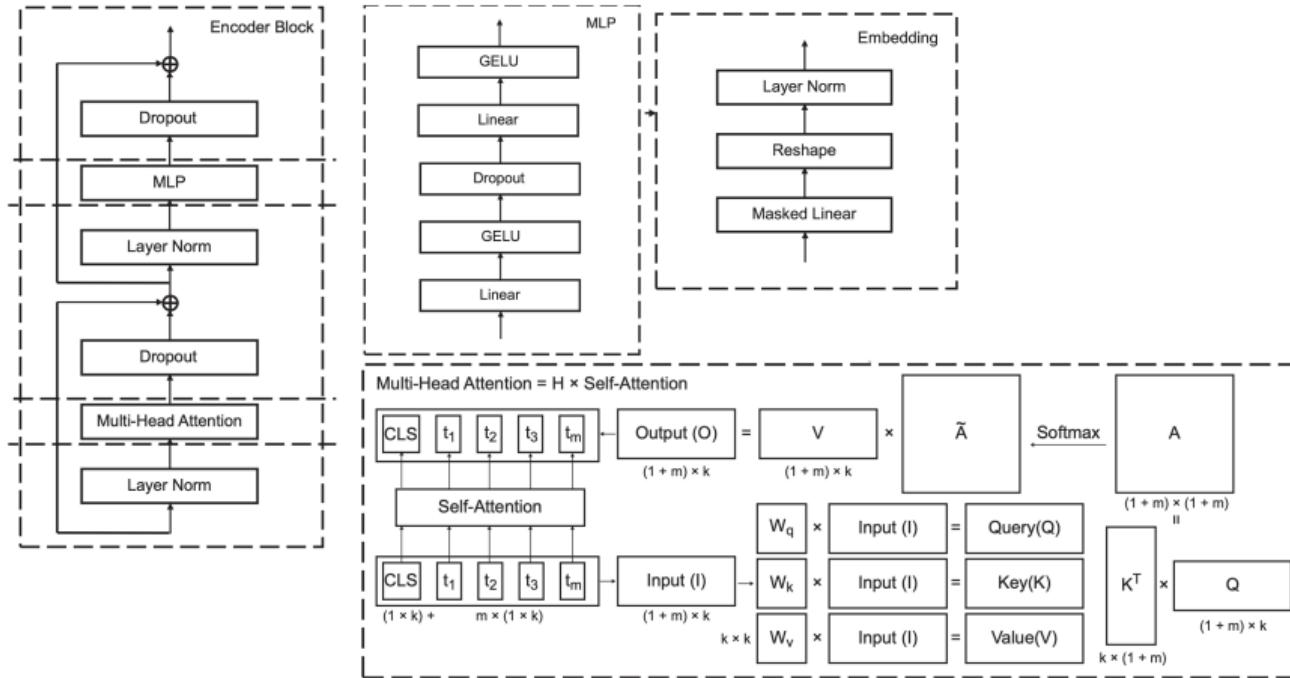
---

The main structure is as follows:



# Model Architecture II

The detailed implementation is as follows:



# Training Details

---

- **Data splitting:**
  - Sample splitting is performed using different studies and biological states (cross-dataset/batch splitting, rather than simple splitting from the same source).
  - 30% of the training set is further divided into a validation set.
- **Loss function and optimizer:**
  - Cross Entropy loss function is used.
  - SGD optimizer is used.
  - Cosine learning rate decay is introduced for the learning rate.
- **Training period:** Convergence within 20 epochs.

# A foundation model of transcription across human cell types

Fu X, Mo S, Buendia A, Laurent AP, Shao A, Alvarez-Torres MDM, Yu T, Tan J, Su J, Sagatelian R, Ferrando AA, Ciccia A, Lan Y, Owens DM, Palomero T, Xing EP, Rabadan R

*Nature*, 2025

DOI: 10.1038/s41586-024-08391-z

<https://www.nature.com/articles/s41586-024-08391-z.pdf>

# Introduction

---

In general, GET is a probabilistic model aiming to predict  $p(E|X, C)$  where:

- $E$ : Gene expression levels (response variable)
- $X$ : Regulatory feature matrix (chromatin accessibility + transcription factor binding sites)
- $C$ : Cell type conditions (even many other conditions, e.g. different experiment methods, sequencing platforms, physiological states, etc.)

It adopts a two-stage framework: pre-training and fine-tuning.

- **Pre-training:** In pretraining, it learns the distribution pattern of regulatory features given the cell types  $p(X|C)$  with unsupervised learning.
- **Fine-tuning:** Adapting the pre-trained model to specific cell types to predict  $f \circ p : X \rightarrow E$ .

The ultimate goal is to generalize the model's predictions across different cell types and conditions:

$$\min_{\theta} \mathbb{E}_{C^{\text{test}}} [\mathcal{L}(f_{\theta}(X^{\text{test}}), E^{\text{test}})]$$

## Notation I

---

- For a gene  $g$  in cell type  $c$ , its regulatory region window (2 Mbp long) can detect  $N$  regions/peaks, denoted as  $\{r_i\}_{i=1}^N$  (here  $N = 200$ ).
- For each region  $r_i$ :
  - We perform motif analysis to obtain motif scores  $\mathbf{m}_i \in \mathbb{R}^{d_m}$ , where  $d_m = 282$  is the motif dimension.
  - We have chromatin accessibility information  $a_i \in \mathbb{R}$ , measured through scATAC-seq experiments using logCPM counts.
- We concatenate motif scores and accessibility to obtain the feature representation:

$$\mathbf{x}_i = \begin{bmatrix} \mathbf{m}_i \\ a_i \end{bmatrix} \in \mathbb{R}^d$$

where  $d := d_m + 1 = 283$ .

## Notation II

---

- All  $N$  regions are encoded to obtain the input feature matrix:

$$X = \begin{bmatrix} \mathbf{x}_1^\top \\ \mathbf{x}_2^\top \\ \vdots \\ \mathbf{x}_N^\top \end{bmatrix} \in \mathbb{R}^{N \times d}$$

## Key Architecture: Region Embedding

---

For each region  $r_i$ , we first transform its feature vector  $\mathbf{x}_i$  into a higher dimensional space through linear transformation (without activation):

- Original feature matrix:  $X \in \mathbb{R}^{N \times d}$
- Linear transformation:  $W_{\text{Emb}} \in \mathbb{R}^{d \times D}$ , where  $D = 768$  in the paper
- Transformed feature matrix:

$$X' := \text{RegionEmb}(X) = XW_{\text{Emb}} \in \mathbb{R}^{N \times D}$$

Note: Since motifs in the original vectors have high correlation, nonlinear transformations are avoided early to prevent compression/interference.

## Key Architecture: Token-wise Self-Attention I

---

To model regulatory relationships between peaks (including cis-interactions and trans-interactions), we employ a 12-layer Transformer Encoder with standard self-attention mechanism:

- Input: Region embeddings  $X' \in \mathbb{R}^{N \times D}$  from previous layer
- For each attention head:
  - Linear projections to Query, Key, Value spaces:

$$Q = X' W_q \in \mathbb{R}^{N \times d_k}$$

$$K = X' W_k \in \mathbb{R}^{N \times d_k}$$

$$V = X' W_v \in \mathbb{R}^{N \times d_v}$$

where  $W_q, W_k \in \mathbb{R}^{D \times d_k}, W_v \in \mathbb{R}^{D \times d_v}$  are learnable

## Key Architecture: Token-wise Self-Attention II

---

- Compute scaled dot-product attention:

$$O_h = \text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^\top}{\sqrt{d_k}} \right) V$$

- Multi-head processing:
  - Compute multiple attention heads in parallel
  - Concatenate outputs and project through fully connected layer
  - Add residual connections and layer normalization
  - Include feed-forward layers between attention blocks

The 12-layer architecture enables modeling complex interactions between regulatory regions.

## Training Procedure: Pre-training I

---

The pre-training phase employs a **masked autoencoding** strategy, inspired by models like BERT, to learn meaningful representations of genomic regions.

- **Masking:** For each input matrix  $X \in \mathbb{R}^{N \times d}$ , half of the regions are randomly selected for masking. Let  $\mathcal{M}$  be the set of indices for these masked regions, where  $|\mathcal{M}| = N/2$ .
  - The feature vectors  $x_i$  for all masked regions ( $i \in \mathcal{M}$ ) are replaced by a single, learnable mask token vector  $\mathbf{m} \in \mathbb{R}^d$ . This results in a masked input  $X^{\text{masked}}$ .
- **Reconstruction Task:** The model is trained to reconstruct the original feature vectors of the masked regions from the corrupted input.

The training process is as follows:

## Training Procedure: Pre-training II

---

1. The masked input  $X^{\text{masked}}$  is passed through the Region Embedding layer to get token embeddings:

$$H^{\text{masked}} = \text{RegionEmb}(X^{\text{masked}}) \in \mathbb{R}^{N \times D}$$

2. The embeddings are processed by the Transformer Encoder to produce contextualized representations:

$$Z^{\text{masked}} = \text{Transformer}(H^{\text{masked}}) \in \mathbb{R}^{N \times D}$$

3. A linear decoder head uses the output representations  $z_i^{\text{masked}}$  to predict the original features for the masked regions:

$$\hat{x}_i = z_i^{\text{masked}} W_{\text{dec}} \in \mathbb{R}^d, \quad \forall i \in \mathcal{M}$$

## Training Procedure: Pre-training III

---

4. The loss function is the reconstruction error (e.g., Mean Squared Error) between the predicted and original features for the masked regions:

$$\mathcal{L} = \sum_{i \in \mathcal{M}} \|\hat{x}_i - x_i\|_2^2$$

Abstractly, the goal is to train the encoder  $p$  and a prediction head  $g$  to minimize the expected reconstruction error over all possible data samples and masks:

$$\min_{p,g} \mathbb{E}_{X,\mathcal{M}} \left[ \sum_{i \in \mathcal{M}} \|g(p(X^{\text{masked}}))_i - x_i\|^2 \right]$$

## Training Procedure: Pre-training IV

---

The training details in this paper are as follows:

- **Optimizer:** AdamW with a weight decay of 0.05
- **Batch Size:** 256
- **Epochs:** 800, with a 40-epoch linear warmup period
- **Max Learning Rate:**  $1.5 \times 10^{-4}$
- **Compute:** Trained for 7 days on 16 NVIDIA V100 GPUs

## Training Procedure: Fine-tuning I

---

After extensive masked pre-training, we fine-tune the model for gene expression prediction tasks.

- **Task:** Predict the RNA expression level  $\hat{y}$  for a gene  $g$  in cell type  $c$ , given its regulatory region features  $X$ .
- **Process:**
  1. Use pre-trained encoder to obtain latent representations:

$$H = p(X) = \begin{bmatrix} h_1^\top \\ \vdots \\ h_N^\top \end{bmatrix} \in \mathbb{R}^{N \times D}$$

2. Apply attention pooling to aggregate region representations:

$$z = \sum_{i=1}^N \alpha_i h_i \in \mathbb{R}^D, \quad \text{where } \alpha_i = \frac{\exp(w^\top h_i)}{\sum_{j=1}^N \exp(w^\top h_j)}$$

## Training Procedure: Fine-tuning II

---

3. Pass pooled representation through an MLP:

$$\hat{y} = f_\phi(z)$$

- **Optimization objective:**

$$\min_{\phi, \theta} \sum_{(g, c)} \mathcal{L}_{\text{expr}}(f_\phi(p_\theta(X_{g, c})), y_{g, c})$$

where  $\mathcal{L}_{\text{expr}}$  is Poisson Negative Log Likelihood loss for count data.

Note: In implementation, this paper adopts LoRA (Hu et al., 2021) for efficient parameter fine-tuning.

## **VCC Baseline and Cell-eval**

<https://github.com/Arclnstitute/cell-eval>